
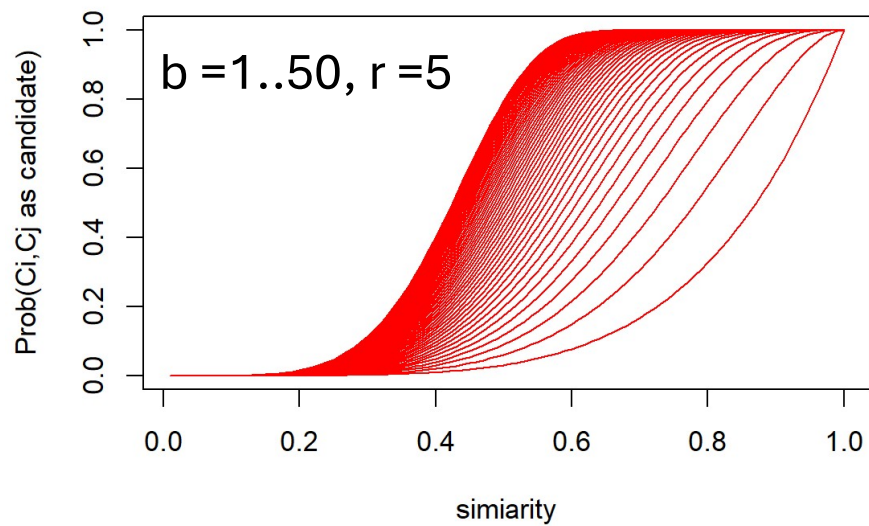




Next step: LSH Tuning

- Tune  $M$ ,  $b$ ,  $r$  to get almost all pairs with similar signatures, but eliminate most pairs that do not have similar signatures

# The S function



LSH: families of functions

# Locality-Sensitive functions

- The example we have previously seen is an example of **family of function** that can be combined to distinguish between pair at a low distance from pairs at a high distance
- Here we will describe other families of functions, beside the minhash one.
- There are **3 conditions** that we need to ensure:
  - **Close pair must be more likely candidates than distant ones**
  - **Statistically independent**: estimate probability that two or more function will give a certain response
  - **Efficient**:
    - Time needed to identify candidate pair much less than computing distance between all pairs
    - Combinable to build functions that are better at avoiding false positive and negatives

# Locality-Sensitive functions

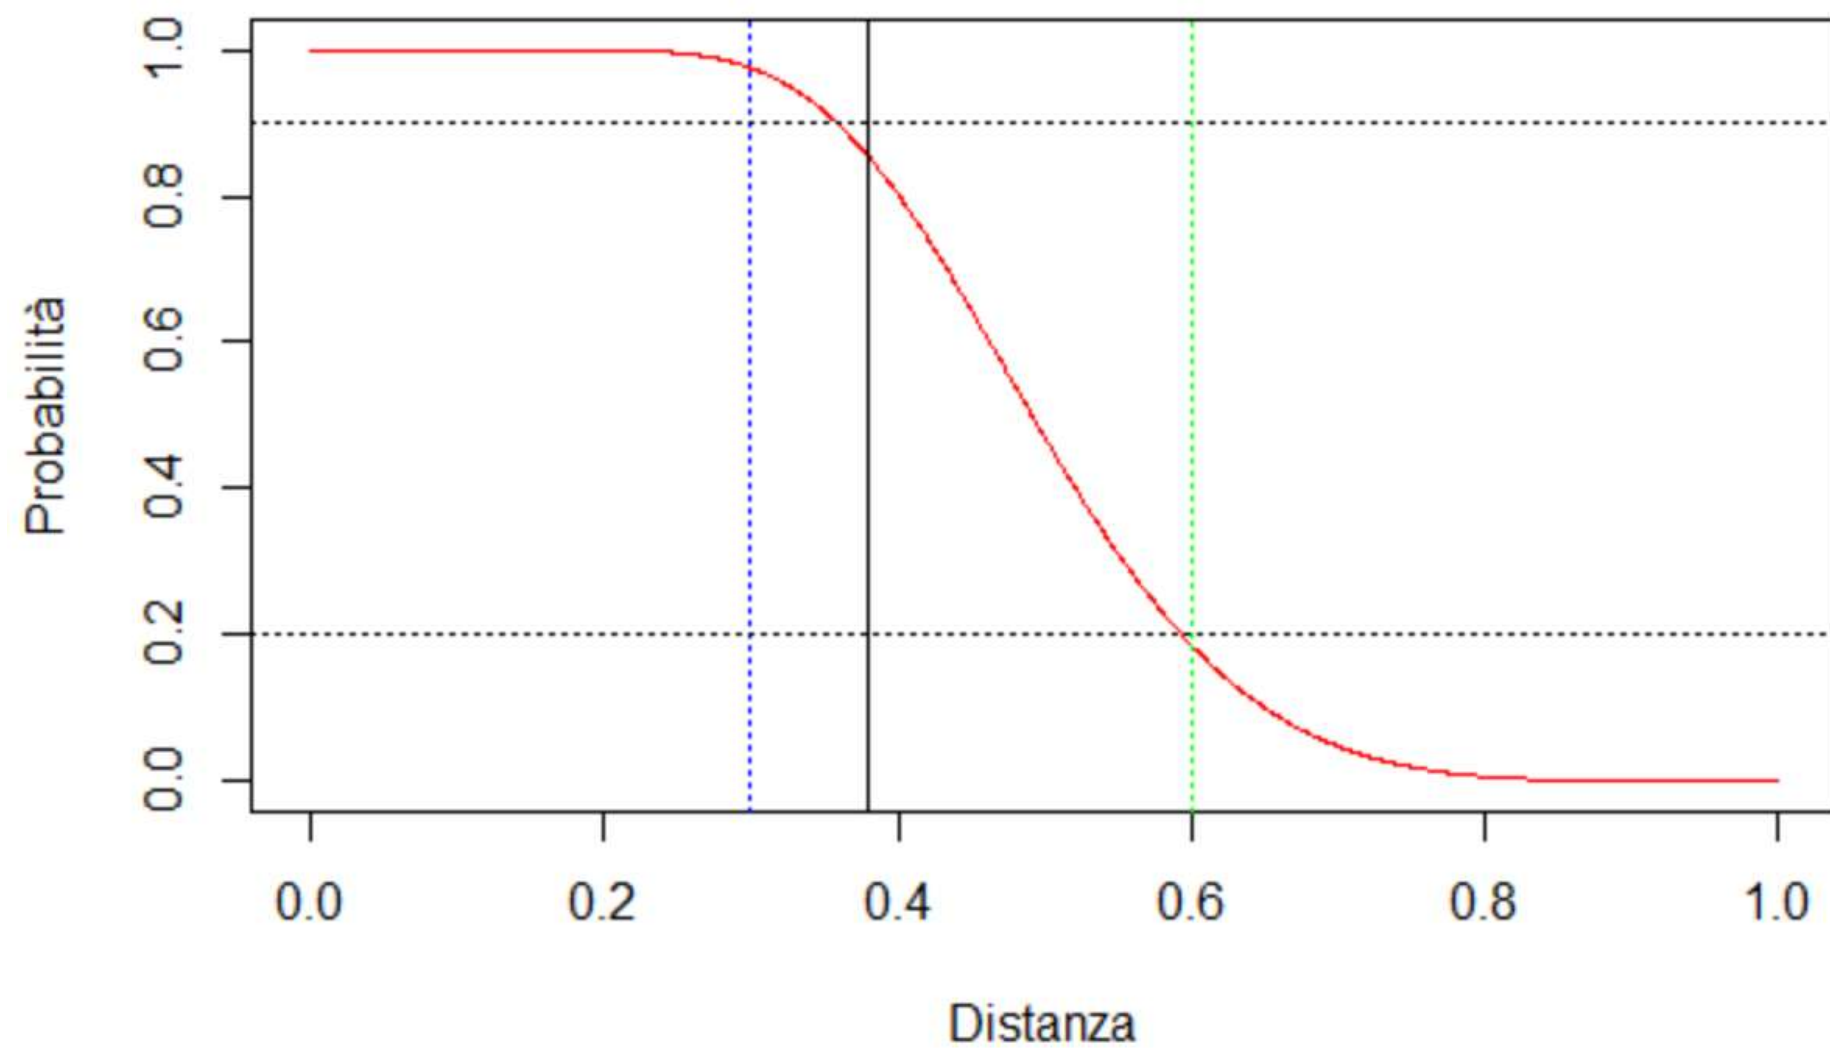
- What is a Locality-Sensitive function?
  - It is a function that takes in input two items and render a decision about whether these items should be a candidate pair.
  - Shorthands:
    - $f(x) = f(y)$  x and y **are** candidate pair
    - $f(x) \neq f(y)$  x and y **are not** candidate pair
- A collection of such functions is called a **family of Locality-Sensitive functions**
- We say that a family **F** is  $(d_1, d_2, p_1, p_2) - sensitive$  if:
  - $\forall f \in F$ :
    - If  $d(x, y) \leq d_1$  the probability that  $f(x) = f(y)$  is at least  $p_1$
    - If  $d(x, y) \geq d_2$  the probability that  $f(x) = f(y)$  is at most  $p_2$

# LSH for min-hashing

- We have
  - $\mathcal{S}$  = space of all sets
  - $d$  = Jaccard distance,
  - $\mathcal{H}$  is family of Min-Hash functions for all permutations of rows
- For any hash function  $h \in \mathcal{H}$ :
$$P[h(x) = h(y)] = 1 - d(x, y)$$
- Simply restates theorem about Min-Hashing in terms of distances rather than similarities


# Locality-Sensitive functions

- The family of **minhash functions**, assuming that we are using a Jaccard Distance, is  
 $(d_1, d_2, 1 - d_1, 1 - d_2) - \text{sensitive}$
- For any  $d_1$  and  $d_2$  such that  $0 < d_1 < d_2 \leq 1$
- For other distances, there is **no guarantee** that it might have a locality-sensitive family of hash functions.
- Min-hash  $H$  is a  $(1/3, 2/3, 2/3, 1/3)$ -sensitive family for  $S$  and  $d$ .





- Can we reproduce the “S-curve” effect we saw before for any LS family?



- The “bands” technique we learned for signature matrices carries over to this more general setting
- Can do LSH with any  $(d1, d2, p1, p2)$ -sensitive family!
- Two constructions:
  - AND construction like “rows in a band”
  - OR construction like “many bands”

# Amplifying Locality-Sensitive functions

- Lets build a new family  $F'$  where each  $f \in F'$  is built from  $r$  members of a family  $F$ , for some fixed  $r$ .
  - $f$  is built from  $\{f_1, \dots, f_r\}$  such that  $f(x) = f(y)$  if and only if  $f_i(x) = f_i(y)$  for all  $i = 1, 2, \dots, r$
- If  $F$  is  $(d_1, d_2, p_1, p_2)$  – *sensitive*  
 $F'$  will be  $(d_1, d_2, (p_1)^r, (p_2)^r)$  – *sensitive*
- We are lowering all probabilities if we choose  $F$  and  $r$  judiciously
- This process is called **AND-construction**


# Hash function independencies

- Independence of hash functions (HFs) really means that the prob. of two HFs saying “yes” is the product of each saying “yes” But two particular hash functions could be highly correlated
- For example, in Min-Hash if their permutations agree in the first one million entries
- However, the probabilities in definition of a LSH-family are over all possible members of  $H$ ,  $H'$  (i.e., average case and not the worst case)


# Amplifying Locality-Sensitive functions

- Now let's build a family  $F'$  where each  $f \in F'$  is built from  $b$  members of a family  $F$ , for some fixed  $b$ .
  - $f$  is built from  $\{f_1, \dots, f_b\}$  such that  $f(x) = f(y)$  if and only if  $f_i(x) = f_i(y)$  for one or more values  $i$
- If  $F$  is  $(d_1, d_2, p_1, p_2) - \text{sensitive}$   
 $F'$  will be  $(d_1, d_2, 1 - (1 - p_1)^b, 1 - (1 - p_2)^b) - \text{sensitive}$
- We are increasing all probabilities if we choose  $F$  and  $b$  judiciously
- This process is called **OR-construction**

- AND and OR constructions can be **cascaded** in order to make the low probability close to 0 and the high probability close to 1.



AND



OR

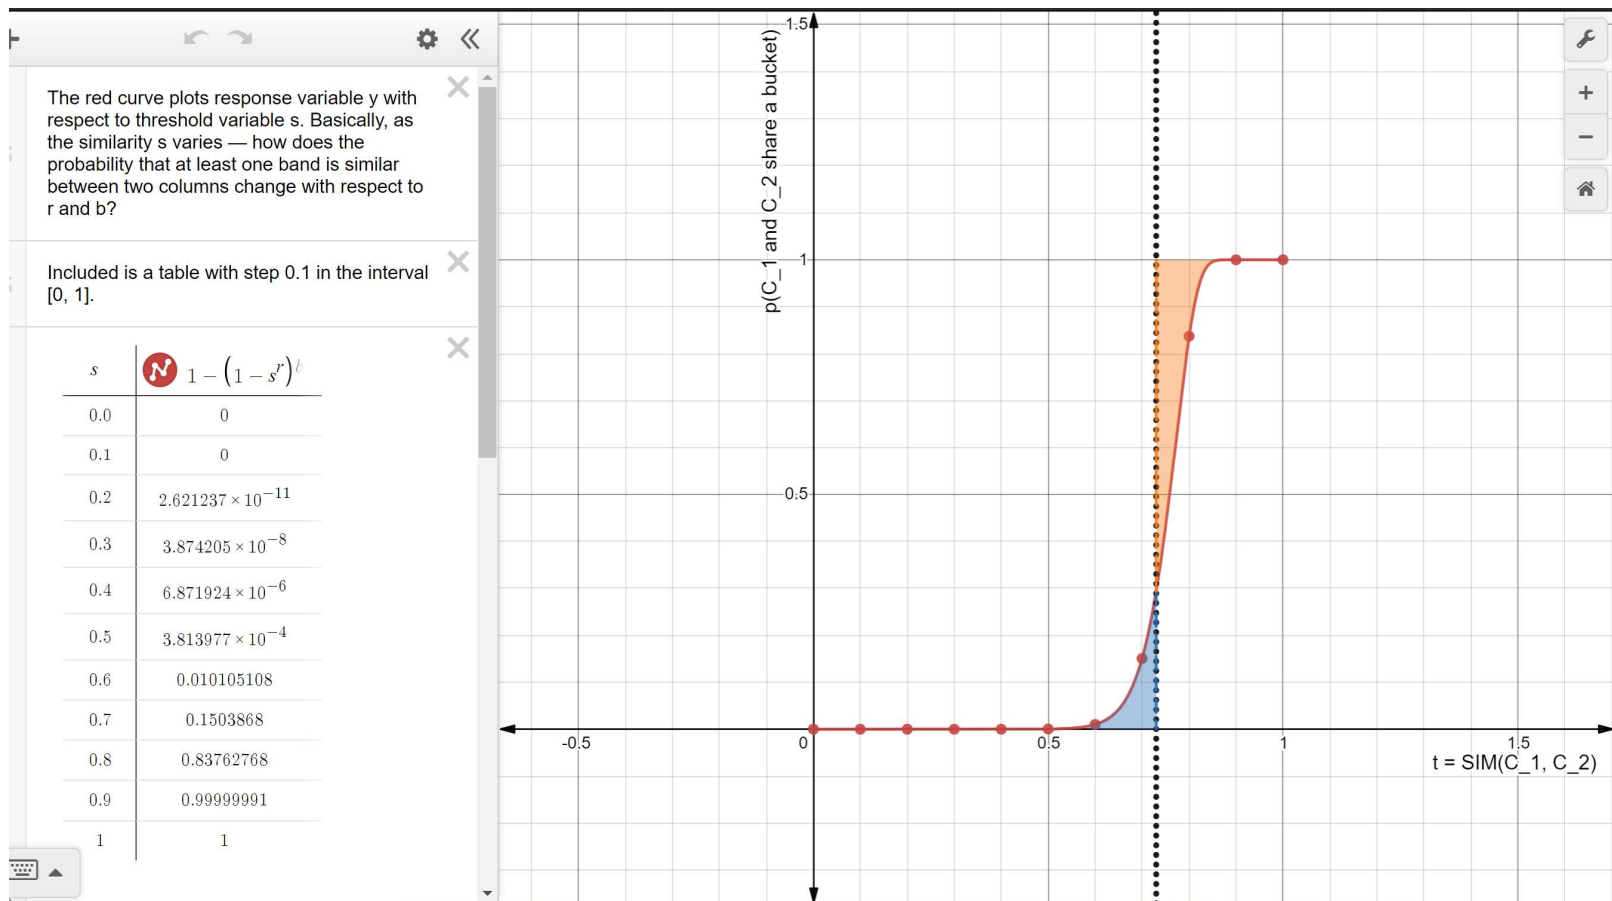
- By choosing  $b$  and  $r$  properly we can make the lower probability close to 0 and the higher close to 1.
- As for the signature matrix, we can use the AND construction followed by the OR construction Or vice-versa
- Or any sequence of AND's and OR's alternating

# Example

- $(0.2, 0.8, 0.8, 0.2)$ -sensitive
- $B=4$   $r = 4$ .
  - AND OR  $(0.2, 0.8, 0.878, 0.0064)$ -sensitive
  - OR AND  $(0.2, 0.8, 0.9936, 0.1215)$ -sensitive
  - Apply again OR and next AND
    - $(0.2, 0.8, 0.9999996, 0.0008715)$ -sensitive

# Demo


- <https://www.desmos.com/calculator/lzzvfjiujn>





- For each AND-OR S-curve  $1-(1-s)^b$ , there is a *threshold*  $t$ , for which  $1-(1-t)^b = t$
- Above  $t$ , high probabilities are increased; below  $t$ , low probabilities are decreased
- You improve the sensitivity as long as the low probability is less than  $t$ , and the high probability is greater than  $t$
- Iterate as you like (computation is not an issue)

- RED curve AND-OR
- GREEN curve AND-OR + AND-OR
- BLUE curve AND-OR + AND-OR + AND-OR
- Threshold 0.73



- Similar observation for the OR-AND type of S-curve:  $(1-(1-s)^b)^r$

# Distance Measures

# Until now: Jaccard Distance

- We talked about **Jaccard Distance and Hamming**
  - measures how close are two sets
- There are other relevant functions to measure the closeness (distance)
- Lets recall some **properties of distance metrics**:
  - $d(x, y) \geq 0$  (no negative distances are allowed)
  - $d(x, y) = 0 \Leftrightarrow x = y$  (distances are positive, except a point from itself)
  - $d(x, y) = d(y, x)$  (distances are symmetric)
  - $d(x, y) \leq d(x, z) + d(z, y)$  (triangle inequality)

# Hamming Distance

- Used with spaces of any vector
- It is computed as the number of components in which two vectors differ.
- **Example:**
  - $x = 10101$
  - $y = 11110$
  - $d(x,y) = 3$

# Hamming Distance

- Building an LSH family of function for **Hamming Distance** is quite simple
- Let  $x$  and  $y$  be two  **$d$ -dimensional vectors**
  - ...we denote with  $h(x, y)$  their hamming distance
- We define a function  $f_i(x)$  to be the  $i$ -th position of vector  $x$ 
  - $f_i(x) = f_i(y)$  if and only if  $x$  and  $y$  agree in the  $i$ -th position.
  - the **probability of  $f(x) = f(y)$**  can be computed as  $1 - h(x, y)/d$
- The **family  $F$**  consisting of the functions previously defined is  
 $(d_1, d_2, 1 - d_1/d, 1 - d_2/d) - \text{sensitive}$
- with  $d_1 < d_2$

# Cosine Distance

- It is a measure defined in **Euclidean spaces** and **discrete Euclidean spaces**.
- It is the cosine of the **angle** that the vectors associated to those points make.
- $d(A, B) = \theta = \arccos(A \cdot B / \|A\| \cdot \|B\|)$
- Has range  **$[0, \pi]$**
- The angle is always in the **range 0 to 180 degrees**.

$$d([x_1, \dots, x_n], [y_1, \dots, y_n]) = 1 - \frac{A \cdot B}{\|A\| \|B\|}$$
$$A \cdot B = \|A\| \|B\| \cos \theta$$

# Cosine Distance

- What is **Cosine Distance** between  $x$  and  $y$ ?
  - the cosine of the angle between them
- For cosine distance, there is a technique
  - called Random Hyperplanes
  - Technique similar to Min-Hashing
- Lets pick a **random hyperplane** and its **normal vector**
  - ...then compute the dot products  $v \cdot x$  and  $v \cdot y$ 
    - two cases:  $\text{sign}(v \cdot x) = \text{sign}(v \cdot y)$  or  $\text{sign}(v \cdot x) \neq \text{sign}(v \cdot y)$


# Cosine Distance

- We define a function  $f(x)$  by choosing a random vector  $v_f$ 
  - $f(x) = f(y)$  if and only  $\text{sign}(v_f \cdot x) = \text{sign}(v_f \cdot y)$ .
- The **family F** consisting of the functions previously defined is
$$(d_1, d_2, (180 - d_1)/180, (180 - d_2)/180)$$
  
– *sensitive*
- with  $d_1 < d_2$



- Each vector  $\mathbf{v}$  determines a hash function  $h_v$  with **two buckets**
- $h_v(x) = +1$  if  $v \cdot x \geq 0$
- $h_v(x) = -1$  if  $v \cdot x < 0$
- LS-family  $\mathbf{H}$  = set of all functions derived from any vector
- **Claim:**
  - For points  $\mathbf{x}$  and  $\mathbf{y}$ ,  $\Pr[h(\mathbf{x}) = h(\mathbf{y})] = 1 - d(\mathbf{x}, \mathbf{y}) / \pi$


# Random hyperplanes



Consider points  $\mathbf{x}$  and  $\mathbf{y}$


■ Let's analyze hyperplanes  $\mathbf{v}$  and  $\mathbf{v}'$

Look in the  
plane of  $\mathbf{x}$   
and  $\mathbf{y}$ .



- $h_v(\mathbf{x}) = +1$  if  $\mathbf{v} \cdot \mathbf{x} \geq 0$
- $h_v(\mathbf{x}) = -1$  if  $\mathbf{v} \cdot \mathbf{x} \leq 0$

- Prob. Red case  $= \theta / \pi$
- Prob green case  $1 - \theta / \pi$



# How to build the signature

- Pick some number of random vectors, and hash your data for each vector
- The result is a **signature** (*sketch*) of **+1**'s and **-1**'s for each data point
- Can be used for LSH like we used the Min-Hash signatures for Jaccard distance
- Amplify using **AND/OR** constructions
- We can demonstrate that by restricting our choice to vectors whose components are **+1** and **-1**, we don't lose randomness.

# Distance in Euclidean spaces

- **n-dimensional Euclidean space:**
  - points are vectors of **n** real numbers.
- For any constant **r**, an  **$L_r$ -norm** is a measure defined as:
- $d([x_1, \dots, x_n], [y_1, \dots, y_n]) = (\sum_{i=1}^n |x_i - y_i|^r)^{1/r}$
- Some famous norms:
  - $L_1$ -norm – **Manhattan Distance**
    - $d([x_1, \dots, x_n], [y_1, \dots, y_n]) = \sum_{i=1}^n |x_i - y_i|$
  - $L_2$ -norm – **Euclidean Distance**
    - $d([x_1, \dots, x_n], [y_1, \dots, y_n]) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$

# Euclidean Distance

- Building an LSH function for Euclidean distance is much more difficult
- 2-dimensional example: let's pick a **random horizontal line**
  - **project** each point into such a line then divide it into **segments of length  $a$** 
    - each segment is a bucket where our **function  $f(x)$**  hashes each point.
- We define **a family  $F$**  by choosing **random lines** through the space and a **bucket size “ $a$ ”** that partitions each line.
  - We will hash points by **projecting** them onto the line  
 **$(d_1, d_2, p_1, p_2) - sensitive$**
- We are not able to determine an expression for  **$p_1$**  and  **$p_2$**  but we can be certain that  **$p_1 > p_2$**  for each  **$d_1 < d_2$**

# Random Projection for Euclidean Distance

- Let  $x$  be a random vector with coordinates selected randomly from a Gaussian Distribution  $N(0,1)$
- Let  $v$  a query point,  $h(v) = x \cdot v$  is the scalar product of the two vectors.
- The scalar projection is then quantized into a set of hash bins with the intention that close items in the original space will fall in the same bin. The hash function is:


$$h^{x,b}(v) = \left\lfloor \frac{x \cdot v + b}{w} \right\rfloor$$

- $w$  is the width of each quantization bin, and  $b$  is a random variable uniformly distributed between 0 and  $w$  that makes the quantization error easier to analyze, with no loss in performance.
- for any points  $p$  and  $q$  in  $\mathbb{R}^d$  that are close to each other, there is a high probability  $\mathbf{p}_1$  that they fall into the same bucket

$$P_h[h(p) = h(q)] \geq p_1 \text{ for } \|p - q\| \leq \mathbf{d}_1$$


- for any points  $p$  and  $q$  in  $\mathbb{R}^d$  that are far apart, there is a low probability  $\mathbf{p}_2 < \mathbf{p}_1$  that they fall into the same bucket

$$P_h[h(p) = h(q)] \leq p_2 \text{ for } \|p - q\| \geq c * d_1 = \mathbf{d}_2$$




■ “Lucky” case:

- Points that are close hash in the same bucket
- Distant points end up in different buckets





■ Two “unlucky” cases:

- **Top:** unlucky quantization
- **Bottom:** unlucky projection







|



# Two cases



- If points are distance  $d < a/2$ , prob. they are in same bucket  $\geq 1 - d/a = 1/2$
- If points are distance  $d > 2a$  apart, then they can be in the same bucket only if  $d \cos \theta \leq a$
- $\cos \theta \leq 1/2$
- $60^\circ < \theta < 90^\circ$ , i.e., at most  $1/3$  probability
- Yields a  $(a/2, 2a, 1/2, 1/3)$ -sensitive family of hash functions for any  $a$

# Edit distance

- Used with **strings**
- The distance between two strings **x** and **y** is the **smallest** number of **insertions** and **deletions** of single characters that will convert x to y.
- **Example:**
  - $x = abcde$
  - $y = acfdeg$
  - $d(x,y) = 3$ 
    - delete «b»
    - insert «f» after «c»
    - insert «g» after «e»

# Other applications of LSH

- **Fingerprint matching:** comparing fingerprints
- **Newspaper articles matching:** distinguishing newspaper article from all the extraneous material
- **Entity resolution:** mapping data records that refer to the same entity (i.e. person)

# Summary: 3 Steps

- **Shingling:** Convert documents to sets
  - We used hashing to assign each shingle an ID
- **Min-Hashing:** Convert large sets to short signatures, while preserving similarity
  - We used **similarity preserving hashing** to generate signatures with property  $\Pr[h_{\pi}(C_1) = h_{\pi}(C_2)] = \text{sim}(C_1, C_2)$
  - We used hashing to get around generating random permutations
- **Locality-Sensitive Hashing:** Focus on pairs of signatures likely to be from similar documents
  - We used hashing to find **candidate pairs** of similarity  $\geq s$

# Notebook

- [Spark Implementation MinHashLSH](#)

<https://colab.research.google.com/drive/1FTVN0dvm-eCGSEIF0d7x3i-PH7BM8yJb?usp=sharing>